

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np # Added for numpy operations
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression # Added for Logistic Regression
from sklearn.neighbors import KNeighborsClassifier # Added for KNN Classifier
from sklearn.metrics import confusion_matrix # Added for confusion_matrix

accuracies = {} # Initialized accuracies dictionary
```

```
df = pd.read_csv('/content/heart.csv')
```

```
df.head()
```

	age	sex	cp	trestbps	chol	fbst	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

Next steps: [Generate code with df](#) [New interactive sheet](#)

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
 --- 
 0   age         303 non-null    int64  
 1   sex         303 non-null    int64  
 2   cp          303 non-null    int64  
 3   trestbps   303 non-null    int64  
 4   chol        303 non-null    int64  
 5   fbs         303 non-null    int64  
 6   restecg    303 non-null    int64  
 7   thalach    303 non-null    int64  
 8   exang       303 non-null    int64  
 9   oldpeak    303 non-null    float64 
 10  slope       303 non-null    int64  
 11  ca          303 non-null    int64  
 12  thal        303 non-null    int64  
 13  target      303 non-null    int64  
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

```
df.shape
```

```
(303, 14)
```

```
df.isnull().sum()
```

```
0
age 0
sex 0
cp 0
trestbps 0
chol 0
fbs 0
restecg 0
thalach 0
exang 0
oldpeak 0
slope 0
ca 0
thal 0
target 0
```

dtype: int64

```
df.duplicated().sum()
```

```
np.int64(1)
```

```
df.drop_duplicates(inplace=True, keep = 'first')
```

```
df.shape
```

```
(302, 14)
```

```
df.describe().T
```

	count	mean	std	min	25%	50%	75%	max	grid	hist
age	302.0	54.420530	9.047970	29.0	48.00	55.5	61.00	77.0		
sex	302.0	0.682119	0.466426	0.0	0.00	1.0	1.00	1.0		
cp	302.0	0.963576	1.032044	0.0	0.00	1.0	2.00	3.0		
trestbps	302.0	131.602649	17.563394	94.0	120.00	130.0	140.00	200.0		
chol	302.0	246.500000	51.753489	126.0	211.00	240.5	274.75	564.0		
fbs	302.0	0.149007	0.356686	0.0	0.00	0.0	0.00	1.0		
restecg	302.0	0.526490	0.526027	0.0	0.00	1.0	1.00	2.0		
thalach	302.0	149.569536	22.903527	71.0	133.25	152.5	166.00	202.0		
exang	302.0	0.327815	0.470196	0.0	0.00	0.0	1.00	1.0		
oldpeak	302.0	1.043046	1.161452	0.0	0.00	0.8	1.60	6.2		
slope	302.0	1.397351	0.616274	0.0	1.00	1.0	2.00	2.0		
ca	302.0	0.718543	1.006748	0.0	0.00	0.0	1.00	4.0		
thal	302.0	2.314570	0.613026	0.0	2.00	2.0	3.00	3.0		
target	302.0	0.543046	0.498970	0.0	0.00	1.0	1.00	1.0		

```
df['target'].value_counts()
```

target	count
1	164
0	138

dtype: int64

```
df = df[df['ca'] < 4] #CA IN KAGGLE 0<3 delete any num bigger than 3
DF = df[df['thal'] > 0] #THAL IN KAGGLE delete any num less than 0
print(df.shape)
print(f'len of data: {len(df)}')
```

```
(298, 14)
len of data: 298
```

```
df.columns
```

```
Index(['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg', 'thalach',
       'exang', 'oldpeak', 'slope', 'ca', 'thal', 'target'],
      dtype='object')
```

```
df = df.rename(
    columns = {
        'cp' : 'chest_pain_type',
        'trestbps': 'resting_blood_pressure',
        'chol': 'cholesterol',
        'fbs': 'fasting_blood_sugar',
        'restecg' : 'resting_electrocardiogram',
        'thalach': 'max_heart_rate_achieved',
        'exang': 'exercise_induced_angina',
        'oldpeak': 'st_depression',
        'slope': 'st_slope',
        'ca': 'num_major_vessels',
        'thal': 'thalassemia'},
    errors = "raise")
```

```
df.head()
```

	age	sex	chest_pain_type	resting_blood_pressure	cholesterol	fasting_blood_sugar	resting_electrocardiogram	max_heart
0	63	1		3	145	233	1	0
1	37	1		2	130	250	0	1
2	41	0		1	130	204	0	0
3	56	1		1	120	236	0	1
4	57	0		0	120	354	0	1

Next steps: [Generate code with df](#) [New interactive sheet](#)

```
df['sex'][df['sex'] == 0] = 'female'
df['sex'][df['sex'] == 1] = 'male'
```

/tmp/ipython-input-3284524103.py:1: FutureWarning: ChainedAssignmentError: behaviour will change in pandas 3.0!
You are setting values through chained assignment. Currently this works in certain cases, but when using Copy-on-Write (which is the default behaviour), it can lead to unexpected results. A typical example is when you are setting values in a column of a DataFrame, like:

```
df["col"][row_indexer] = value
```

Use `df.loc[row_indexer, "col"] = values` instead, to perform the assignment in a single step and ensure this keeps updating.

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view

```
df['sex'][df['sex'] == 0] = 'female'
```

/tmp/ipython-input-3284524103.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view

```
df['sex'][df['sex'] == 0] = 'female'
```

/tmp/ipython-input-3284524103.py:1: FutureWarning: Setting an item of incompatible dtype is deprecated and will raise an error.

```
df['sex'][df['sex'] == 0] = 'female'
```

```
df.dtypes
```

```
0
age          int64
sex          object
chest_pain_type    int64
resting_blood_pressure  int64
cholesterol      int64
fasting_blood_sugar  int64
resting_electrocardiogram  int64
max_heart_rate_achieved  int64
exercise_induced_angina  int64
st_depression     float64
st_slope          int64
num_major_vessels int64
thalassemia       int64
target           int64

dtype: object
```

```
count = 0
for i in df.dtypes:
    if i == 'object':
        count += 1

print(count)
```

```
1
```

```
df.describe(include='object')
```

sex	grid icon
count	298
unique	2
top	male
freq	202

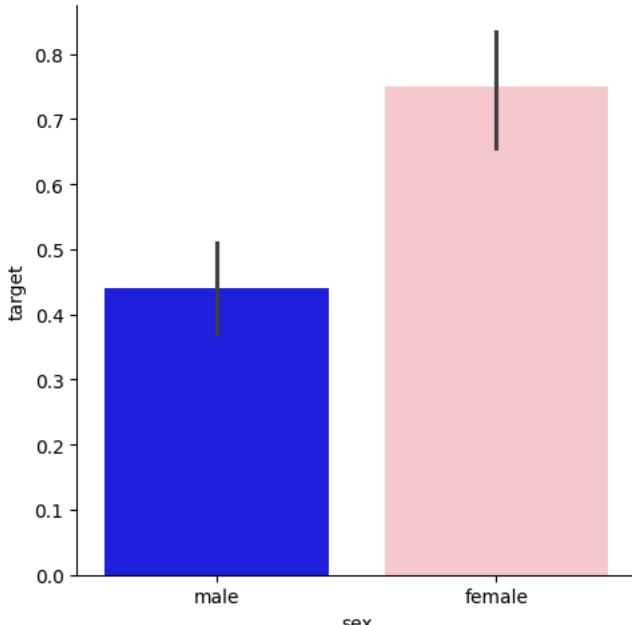
```
Start coding or generate with AI.
```

```
get_object_columns = [col for col in df.columns if df[col].dtypes == 'object']
get_object_columns
```

```
['sex']
```

```
plt.figure(figsize = (16,10))
for i in get_object_columns:
    sns.catplot(y = "target", x = i, hue = "sex", data = df, kind = "bar",
                 palette={"male": "blue", "female": "pink"})
```

<Figure size 1600x1000 with 0 Axes>



```
def label_encode_cat_features(data, cat_features):
    label_encoder = LabelEncoder()
    data_encoded = data.copy()

    for col in cat_features:
        data_encoded[col] = label_encoder.fit_transform(data[col])
    data = data_encoded
    return data
```

```
cat_features = ['sex', 'fasting_blood_sugar', 'exercise_induced_angina', 'target',
                'chest_pain_type', 'resting_electrocardiogram', 'st_slope', 'thalassemia']
```

```
df = label_encode_cat_features(df, cat_features)
```

```
df.head()
```

	age	sex	chest_pain_type	resting_blood_pressure	cholesterol	fasting_blood_sugar	resting_electrocardiogram	max_heart
0	63	1		3	145	233	1	0
1	37	1		2	130	250	0	1
2	41	0		1	130	204	0	0
3	56	1		1	120	236	0	1
4	57	0		0	120	354	0	1

Next steps: [Generate code with df](#) [New interactive sheet](#)

```
X = df.iloc[ : , :-1]
y = df.iloc[ : , -1]
```

```
X.head(2)
```

	age	sex	chest_pain_type	resting_blood_pressure	cholesterol	fasting_blood_sugar	resting_electrocardiogram	max_heart
0	63	1		3	145	233	1	0
1	37	1		2	130	250	0	1

Next steps: [Generate code with X](#) [New interactive sheet](#)

```
y.head(2)
```

```
target
_____
0      1
1      1
dtype: int64
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=42, shuffle = True)
```

```
lr = LogisticRegression()
lr.fit(X_train,y_train)
acc = lr.score(X_test,y_test)*100

accuracies['Logistic Regression'] = acc
print("Test Accuracy {:.2f}%".format(acc))
```

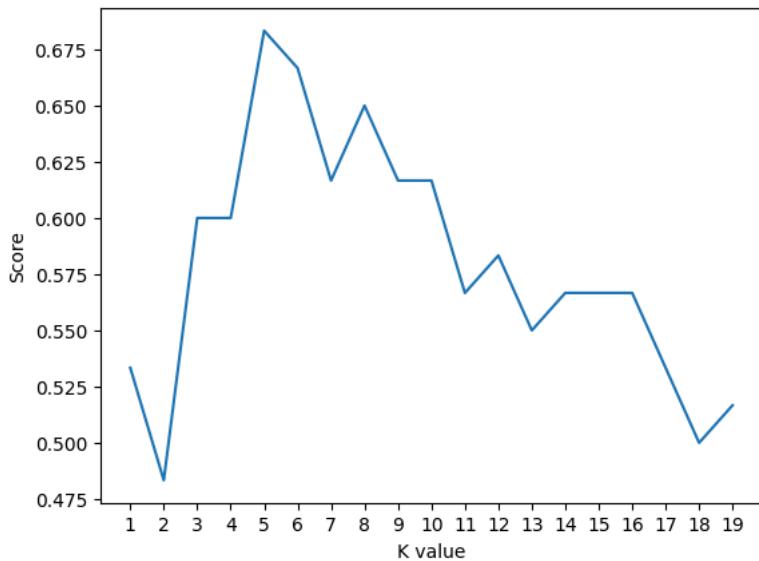
Test Accuracy 81.67%
/usr/local/lib/python3.12/dist-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result()

```
scoreList = []
for i in range(1,20):
    knn = KNeighborsClassifier(n_neighbors = i) # n_neighbors means k
    knn.fit(X_train, y_train)
    scoreList.append(knn.score(X_test, y_test))

plt.plot(range(1,20), scoreList)
plt.xticks(np.arange(1,20,1))
plt.xlabel("K value")
plt.ylabel("Score")
plt.show()

acc = max(scoreList)*100
accuracies['KNN'] = acc
print("Maximum KNN Score is {:.2f}%".format(acc))
```



Maximum KNN Score is 68.33%

```
accuracies
{'Logistic Regression': 81.66666666666667, 'KNN': 68.33333333333333}
```

```
from sklearn.svm import SVC  
  
svc = SVC()  
svc.fit(X_train, y_train)  
  
acc = svc.score(X_test,y_test)*100  
accuracies['SVC'] = acc  
print(f"Test Accuracy of SVC Algorithm: {acc:.2f}%")
```

Test Accuracy of SVC Algorithm: 65.00%

```
from sklearn.naive_bayes import GaussianNB  
nb = GaussianNB()  
nb.fit(X_train, y_train)  
  
acc = nb.score(X_test,y_test)*100  
accuracies['Naive Bayes'] = acc  
print(f"Accuracy of Naive Bayes: {acc:.2f}%")
```

Accuracy of Naive Bayes: 83.33%

```
from sklearn.tree import DecisionTreeClassifier  
dtc = DecisionTreeClassifier()  
dtc.fit(X_train, y_train)  
  
acc = dtc.score(X_test, y_test)*100  
accuracies['Decision Tree'] = acc  
print(f"Decision Tree Test Accuracy {acc:.2f}%")
```

Decision Tree Test Accuracy 70.00%

```
from sklearn.ensemble import RandomForestClassifier  
rf = RandomForestClassifier(n_estimators = 1000, random_state = 1)  
rf.fit(X_train, y_train)  
  
acc = rf.score(X_test,y_test)*100  
accuracies['Random Forest'] = acc  
print(f"Random Forest Algorithm Accuracy Score : {acc:.2f}%")
```

Random Forest Algorithm Accuracy Score : 85.00%

```
colors = ["skyblue", "salmon", "lightgreen", "gold", "lightcoral", "plum"]  
  
sns.set_style("whitegrid")  
plt.figure(figsize=(16, 5))  
plt.yticks(np.arange(0, 110, 10))  
plt.ylabel("Accuracy %")  
plt.xlabel("Algorithms")  
  
ax = sns.barplot(x=list(accuracies.keys()), y=list(accuracies.values()), palette=colors)  
  
for i, v in enumerate(accuracies.values()):  
    ax.text(i, v + 1, f"{v:.2f}%", color='black', ha='center', fontweight='bold')  
  
plt.title("Model Accuracy Comparison")  
plt.ylim(0, 110)  
plt.show()
```

```
/tmp/ipython-input-2054633348.py:10: FutureWarning:  
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and  
ax = sns.barplot(x=list(accuracies.keys()), y=list(accuracies.values()), palette=colors)
```

Model Accuracy Comparison

```
y_pred_lr = lr.predict(X_test)  
y_pred_knn = knn.predict(X_test)  
y_pred_svc = svc.predict(X_test)  
y_pred_nb = nb.predict(X_test)  
y_pred_dtc = dtc.predict(X_test)  
y_pred_rf = rf.predict(X_test)
```



```
cm_lr = confusion_matrix(y_test,y_pred_lr)  
cm_knn = confusion_matrix(y_test,y_pred_knn)  
cm_svc = confusion_matrix(y_test,y_pred_svc)  
cm_nb = confusion_matrix(y_test,y_pred_nb)  
cm_dtc = confusion_matrix(y_test,y_pred_dtc)  
cm_rf = confusion_matrix(y_test,y_pred_rf)
```

Algorithms

```
models = [  
    ("Logistic Regression", cm_lr),  
    ("K Nearest Neighbors", cm_knn),  
    ("Support Vector Machine", cm_svc),  
    ("Naive Bayes", cm_nb),  
    ("Decision Tree Classifier", cm_dtc),  
    ("Random Forest", cm_rf)  
]  
  
plt.figure(figsize=(24, 12))  
plt.suptitle("Confusion Matrixes", fontsize=24)  
plt.subplots_adjust(wspace=0.4, hspace=0.4)  
  
for i, (title, cm) in enumerate(models, 1):  
    plt.subplot(2, 3, i)  
    plt.title(f"{title} Confusion Matrix")  
    sns.heatmap(cm, annot=True, cmap="Blues", fmt="d", cbar=False, annot_kws={"size": 24})  
  
plt.show()
```

Confusion Matrixes

